

1. Upload the Dataset

```
In [1]: from google.colab import files
        uploaded = files.upload()
```

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving Scaled_Customer_Churn_Data.csv to Scaled_Customer_Churn_Data.csv

2. Load the Dataset

```
In [2]: import pandas as pd
        df = pd.read_csv('Scaled_Customer_Churn_Data.csv')
        df.head()
```

```
Out[2]:
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines
0	-1.009430	-0.440327	1.035617	-0.652305	-1.280248	-3.056334	0.06266
1	0.990658	-0.440327	-0.965608	-0.652305	0.064303	0.327189	-0.99156
2	0.990658	-0.440327	-0.965608	-0.652305	-1.239504	0.327189	-0.99156
3	0.990658	-0.440327	-0.965608	-0.652305	0.512486	-3.056334	0.06266
4	-1.009430	-0.440327	-0.965608	-0.652305	-1.239504	0.327189	-0.99156

3. Data Exploration

```
In [3]: df.info()
        df.describe()
        df.columns
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7032 entries, 0 to 7031
Data columns (total 20 columns):
#   Column                Non-Null Count  Dtype
---  -
0   gender                 7032 non-null   float64
1   SeniorCitizen          7032 non-null   float64
2   Partner                7032 non-null   float64
3   Dependents             7032 non-null   float64
4   tenure                 7032 non-null   float64
5   PhoneService           7032 non-null   float64
6   MultipleLines          7032 non-null   float64
7   InternetService        7032 non-null   float64
8   OnlineSecurity         7032 non-null   float64
9   OnlineBackup           7032 non-null   float64
10  DeviceProtection       7032 non-null   float64
11  TechSupport            7032 non-null   float64
12  StreamingTV            7032 non-null   float64
13  StreamingMovies        7032 non-null   float64
14  Contract               7032 non-null   float64
15  PaperlessBilling       7032 non-null   float64
16  PaymentMethod          7032 non-null   float64
17  MonthlyCharges         7032 non-null   float64
18  TotalCharges           7032 non-null   float64
19  Churn_Yes              7032 non-null   int64
dtypes: float64(19), int64(1)
memory usage: 1.1 MB

```

```

Out[3]: Index(['gender', 'SeniorCitizen', 'Partner', 'Dependents', 'tenure',
              'PhoneService', 'MultipleLines', 'InternetService', 'OnlineSecurity',
              'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV',
              'StreamingMovies', 'Contract', 'PaperlessBilling', 'PaymentMethod',
              'MonthlyCharges', 'TotalCharges', 'Churn_Yes'],
              dtype='object')

```

4. Check for Missing Values and Duplicates

```


In [4]: print("Missing values:\n", df.isnull().sum())
        print("Duplicates:", df.duplicated().sum())

```

```
Missing values:
  gender          0
SeniorCitizen    0
Partner          0
Dependents       0
tenure           0
PhoneService     0
MultipleLines    0
InternetService  0
OnlineSecurity   0
OnlineBackup     0
DeviceProtection 0
TechSupport      0
StreamingTV      0
StreamingMovies  0
Contract         0
PaperlessBilling 0
PaymentMethod    0
MonthlyCharges   0
TotalCharges     0
Churn_Yes        0
dtype: int64
Duplicates: 22
```

5. Visualize a Few Features

```
In [5]: import seaborn as sns
import matplotlib.pyplot as plt
sns.countplot(x='Churn_Yes', data=df)
plt.title('Churn Distribution')
plt.xlabel('Churn (0 = No, 1 = Yes)')
plt.ylabel('Count')
plt.show()
```



6. Identify Target and Features

```
In [6]: X = df.drop('Churn_Yes', axis=1)
        y = df['Churn_Yes']
```

7. Convert Categorical Columns to Numerical

```
In [ ]: # Already done in this dataset
```

8. One-Hot Encoding

```
In [7]: # Already done in this dataset
```

9. Feature Scaling

```
In [8]: # Already done in this dataset
```

10. Train-Test Split

```
In [9]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, |
```

11. Model Building

```
In [10]: from sklearn.ensemble import RandomForestClassifier
model = RandomForestClassifier()
model.fit(X_train, y_train)
```

Out[10]:

```
▼ RandomForestClassifier
RandomForestClassifier()
```

12. Evaluation

```
In [11]: from sklearn.metrics import classification_report, confusion_matrix
y_pred = model.predict(X_test)
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

```
[[934  99]
 [195 179]]
```

	precision	recall	f1-score	support
0	0.83	0.90	0.86	1033
1	0.64	0.48	0.55	374
accuracy			0.79	1407
macro avg	0.74	0.69	0.71	1407
weighted avg	0.78	0.79	0.78	1407

13. Make Predictions from New Input

```
In [12]: import numpy as np
sample_input = np.array([X_test.iloc[0]])
model.predict(sample_input)
```

```
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739:
UserWarning: X does not have valid feature names, but RandomForestClassifier
fitted with feature names
warnings.warn(
```

```
Out[12]: array([0])
```

14. Convert to DataFrame and Encode

```
In [13]: # Not required, already encoded
```

15. Predict the Final Grade

```
In [14]: final_prediction = model.predict(sample_input)
print('Churn Prediction:', 'Yes' if final_prediction[0] == 1 else 'No')

Churn Prediction: No
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739:
UserWarning: X does not have valid feature names, but RandomForestClassifier
fitted with feature names
  warnings.warn(
```

16. Deployment - Building an Interactive App

```
In [15]: !pip install gradio
import gradio as gr
```

Collecting gradio
 Downloading gradio-5.29.0-py3-none-any.whl.metadata (16 kB)
 Collecting aiofiles<25.0,>=22.0 (from gradio)
 Downloading aiofiles-24.1.0-py3-none-any.whl.metadata (10 kB)
 Requirement already satisfied: anyio<5.0,>=3.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (4.9.0)
 Collecting fastapi<1.0,>=0.115.2 (from gradio)
 Downloading fastapi-0.115.12-py3-none-any.whl.metadata (27 kB)
 Collecting ffmpeg (from gradio)
 Downloading ffmpeg-0.5.0-py3-none-any.whl.metadata (3.0 kB)
 Collecting gradio-client==1.10.0 (from gradio)
 Downloading gradio_client-1.10.0-py3-none-any.whl.metadata (7.1 kB)
 Collecting groovy~=0.1 (from gradio)
 Downloading groovy-0.1.2-py3-none-any.whl.metadata (6.1 kB)
 Requirement already satisfied: httpx>=0.24.1 in /usr/local/lib/python3.11/dist-packages (from gradio) (0.28.1)
 Requirement already satisfied: huggingface-hub>=0.28.1 in /usr/local/lib/python3.11/dist-packages (from gradio) (0.30.2)
 Requirement already satisfied: jinja2<4.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (3.1.6)
 Requirement already satisfied: markupsafe<4.0,>=2.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (3.0.2)
 Requirement already satisfied: numpy<3.0,>=1.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (2.0.2)
 Requirement already satisfied: orjson~=3.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (3.10.18)
 Requirement already satisfied: packaging in /usr/local/lib/python3.11/dist-packages (from gradio) (24.2)
 Requirement already satisfied: pandas<3.0,>=1.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (2.2.2)
 Requirement already satisfied: pillow<12.0,>=8.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (11.2.1)
 Requirement already satisfied: pydantic<2.12,>=2.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (2.11.4)
 Collecting pydub (from gradio)
 Downloading pydub-0.25.1-py2.py3-none-any.whl.metadata (1.4 kB)
 Collecting python-multipart>=0.0.18 (from gradio)
 Downloading python_multipart-0.0.20-py3-none-any.whl.metadata (1.8 kB)
 Requirement already satisfied: pyyaml<7.0,>=5.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (6.0.2)
 Collecting ruff>=0.9.3 (from gradio)
 Downloading ruff-0.11.8-py3-none-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (25 kB)
 Collecting safehttpx<0.2.0,>=0.1.6 (from gradio)
 Downloading safehttpx-0.1.6-py3-none-any.whl.metadata (4.2 kB)
 Collecting semantic-version~=2.0 (from gradio)
 Downloading semantic_version-2.10.0-py2.py3-none-any.whl.metadata (9.7 kB)
 Collecting starlette<1.0,>=0.40.0 (from gradio)
 Downloading starlette-0.46.2-py3-none-any.whl.metadata (6.2 kB)
 Collecting tomlkit<0.14.0,>=0.12.0 (from gradio)
 Downloading tomlkit-0.13.2-py3-none-any.whl.metadata (2.7 kB)
 Requirement already satisfied: typer<1.0,>=0.12 in /usr/local/lib/python3.11/dist-packages (from gradio) (0.15.3)
 Requirement already satisfied: typing-extensions~=4.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (4.13.2)
 Collecting uvicorn>=0.14.0 (from gradio)
 Downloading uvicorn-0.34.2-py3-none-any.whl.metadata (6.5 kB)
 Requirement already satisfied: fsspec in /usr/local/lib/python3.11/dist-packages (from gradio-client==1.10.0->gradio) (2025.3.2)
 Requirement already satisfied: websockets<16.0,>=10.0 in /usr/local/lib/python3.11/dist-packages (from gradio-client==1.10.0->gradio) (15.0.1)

17. Create a Prediction Function

```
In [16]: def predict_churn(*input_data):  
         input_array = np.array(input_data).reshape(1, -1)  
         prediction = model.predict(input_array)  
         return 'Churn' if prediction[0] == 1 else 'No Churn'
```

18. Create the Gradio Interface

```
In [ ]: inputs = [gr.Number(label=f) for f in X.columns]  
         gr.Interface(fn=predict_churn, inputs=inputs, outputs='text', title='Customer Churn Prediction')  
  
Colab notebook detected. This cell will run indefinitely so that you can see  
and logs. To turn off, set debug=False in launch().  
* Running on public URL: https://6af875ed951eebc0cd.gradio.live  
  
This share link expires in 1 week. For free permanent hosting and GPU upgrade  
'gradio deploy' from the terminal in the working directory to deploy to Hugging  
Face Spaces (https://huggingface.co/spaces)
```

█

19. Student Performance Predictor

```
In [ ]: # This is a churn prediction model; adjust title if necessary
```